

PolicyPulse 360

Customer 360 + Next Best Action for insurance retention teams

Team Neural Mavericks · Snowflake CoCo CLI Hackathon (GCC Edition)

Challenge: Customer 360 and Next Best Action

Snowflake AI Data Cloud · Cortex AI SQL (SENTIMENT · AI_CLASSIFY · AI_COMPLETE · Search) · 3 Cortex Code CLI skills · Streamlit


An agent can't save a customer they can't see

- Indian insurers run policy admin, claims, billing and the call centre as separate systems.
 - A retention agent opens 4+ screens and replays call recordings, taking ~20 min per customer (est.)
- The strongest churn signals are buried in unstructured call transcripts:
 - "I want to port my policy" · "claim pending 45 days" · "an aggregator quoted ₹4,000 less"
- Result: churn is discovered at lapse, offers are generic, and upsell moments are missed.
 - (new baby → health top-up, new home → Home Protect)
- Persona: relationship manager or retention agent at a multi-line insurer (Motor · Health · Life · Home).
- Goal: go from a customer question to an explainable, ready-to-send action in one screen.


Architecture: 3 modular CoCo CLI skills on Snowflake

PolicyPulse 360 — architecture

Structured + unstructured touchpoints → unified 360 → explainable churn risk → next best action, orchestrated by 3 modular Cortex Code CLI skills



Input → Processing → Output, all driven from CoCo CLI



Skills share Snowflake tables as their contract, so each runs alone or chained. The 360 is a view, so AI signals flow in as soon as Skill 2 runs.
Demo: "What should I do for C0004?" → risk 90/100 → senior retention call + WhatsApp message.

Impact: measurable outcomes

20m → 30s

agent prep per customer
(projected: 4 systems → 1 card)

67 / 300

customers flagged High risk
(22% of book, measured)

81%

recall of truly at-risk customers
(61/75; 91% precision)

₹32.9 L

premium in High-risk band
(23% of ₹1.45 Cr, measured)

- 41 High-risk customers renew within 30 days: a ranked call list for this week.
- Saving 20% of High-risk premium = ₹6.6 L per 300 customers, about ₹219 Cr per 1M customers (projected).
- Every recommendation is explainable: score drivers + cited transcript IDs, so it's audit-friendly.
- Upsell signals (new baby, new home) are routed to cross-sell, not missed.

Measured on the synthetic dataset (seed 42, 300 customers, 819 transcripts). Data source for this deck: snowflake-cortex.

Scalability and next steps

- Cost grows with new interactions, not book size:
 - incremental Cortex enrichment; LLM capped to High/Medium-risk + upsell customers
- XS warehouse + auto-suspend 60s handles this demo. Warehouse size is a one-line change for millions.
- Portable skills: drop .cortex/skills into any account, point at real RAW tables, AS_OF_DATE() → CURRENT_DATE()
- DEMO mode (Parquet) lets anyone open the app; LIVE mode queries Snowflake and calls Cortex on demand.

Next steps

- AI_TRANSCRIBE for raw call audio; Hindi and regional-language calls
- Dynamic tables for near-real-time scoring after every call
- Calibrate weights on real lapse outcomes; A/B uplift by action; push NBAs to CRM / WhatsApp Business
- Cortex Analyst semantic model for governed natural-language Q&A