

YOURSTORY



H2S



COCO CLI

HACKATHON

GCC Edition



Team Name : Neural Mavericks

Team Leader Name : Yugraj Prabhakar Mangate

Team Size : 01

Problem Statement : Customer 360 and Next Best Action – PolicyPulse 360


1. Problem Brief

- **Business problem:** Indian insurers run policy admin, claims, billing and the call centre as separate systems, so churn is discovered at lapse and offers are generic.
- **Persona:** retention agent / relationship manager at a multi-line Indian insurer (Motor · Health · Life · Home).
- **Pain point today:** ~20 min per customer across 4+ screens and call recordings. The strongest churn signals ("I'll port my policy", "claim pending 45 days") are buried in unstructured transcripts.
- **How PolicyPulse 360 improves it:** one 360 card that unifies all touchpoints, an explainable churn score, and a next best action with cited call evidence plus a ready-to-send message.
- **Industry context:** high lapse at renewal, aggregator price-shopping, health portability and IRDAI grievance norms make proactive, explainable retention critical.

2. Architecture Diagram

PolicyPulse 360 — architecture

Structured + unstructured touchpoints → unified 360 → explainable churn risk → next best action, orchestrated by 4 modular Cortex Code CLI skills



4 CoCo CLI skills: \$c360-unify → \$interaction-intel → \$next-best-action → \$ask-policypulse
Structured: customers, policies, claims, payments · **Unstructured:** call transcripts

3. Impact Statement

- **Measured on synthetic data (real Cortex run):** 67 of 300 customers (22%) flagged High risk, catching 81% of truly at-risk customers at 91% precision.
- **Measured on synthetic data:** ₹32.9 L of ₹1.45 Cr annual premium (23%) sits in the High-risk band; 41 High-risk renewals are due within 30 days.
- **Projected:** agent prep time ~20 min → under 30 s per customer (4 systems + recordings → one 360 card).
- **Projected:** saving 20% of High-risk premium = ₹6.6 L per 300 customers, about ₹219 Cr per 1M customers.
- **Scalability:** incremental Cortex enrichment and an LLM capped to high-value customers, so cost tracks new calls, not book size. Runs on an XS warehouse with auto-suspend.
- **Beyond the demo:** portable skills on real RAW tables; next up are AI_TRANSCRIBE for call audio, dynamic tables for real-time scoring, and CRM / WhatsApp push.

Additional Slide: Solution Walkthrough

- **Customer 360 + NBA:** explainable risk gauge and drivers, next best action with cited call IDs, a copyable message, and a call timeline with AI summaries.
- **Portfolio:** ₹ premium at risk by product × risk band, intent × sentiment, the 90-day renewal pipeline, and a ranked call list.
- **Ask PolicyPulse:** plain-English questions answered with Cortex AI_COMPLETE text-to-SQL (read-only). How it works: architecture + skills.
- **Cortex used:** SENTIMENT · AI_CLASSIFY · AI_COMPLETE (GPT-4.1 writes the next best action) · Cortex Search · PolicyPulse Copilot agent with cited answers (\$ask-policypulse).

GitHub: <https://github.com/yugrajmangate-dev/neural-mavericks-policypulse360>

Live app (React): <https://yugrajmangate-dev.github.io/neural-mavericks-policypulse360/>

Live app (Streamlit): <https://neural-mavericks-policypulse360.streamlit.app/>

Demo video: <https://drive.google.com/file/d/1sbVZ61HSDpibyNr2ZdpOXsHa40tJKOqP/view>

YOURSTORY



1125

Coco CLI

HACKATHON

GCC Edition

THANK YOU

